

FineTune with LoRA.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Files

- ..
- banking_assistant_lora
- sample_data
- wandb
- banking_dataset.json
- model_comparison.csv
- training_loss.png

!pip install -q transformers peft datasets accelerate bitsandbytes

59.1/59.1 MB 13.8 MB/s eta 0:00:00

CONFIGURATION

```
[2] [✓ 46s]
import json
import torch
from datasets import Dataset
from transformers import AutoModelForSeq2SeqLM, AutoTokenizer, Seq2SeqTrainingArguments, Seq2SeqTrainer, DataCollatorForSeq2Seq
from peft import LoraConfig, get_peft_model, TaskType, PeftModel

MODEL_NAME = "google/flan-t5-base"
DATA_FILE = "./banking_dataset.json"
OUTPUT_DIR = "./banking_assistant_lora"
NUM_EPOCHS = 3
BATCH_SIZE = 8
```

Load Dataset

```
[4] [✓ 0s]
def load_and_format_data(file_path):
    with open(file_path, "r") as f:
        data = json.load(f)

    formatted_data = []
    for item in data:
        formatted_data.append({
            "input_text": f"banking_assistant: {item['instruction']}",
            "target_text": item['response']
        })
    return Dataset.from_list(formatted_data)

dataset = load_and_format_data(DATA_FILE)
```

Disk 72.86 GB available

FineTune with LoRA.ipynb

File Edit View Insert Runtime Tools Help

Commands + Code + Text Run all

Files

- ..
- banking_assistant_lora
- sample_data
- wandb
- banking_dataset.json
- model_comparison.csv
- training_loss.png

Split into train/test/validation

```
[5] [✓ 0s]
dataset = dataset.train_test_split(test_size=0.1)
print(f"Training samples: {len(dataset['train'])}")
print(f"Validation samples: {len(dataset['test'])}")

Training samples: 720
Validation samples: 88
```

Tokenizer

```
[6]
tokenizer = AutoTokenizer.from_pretrained(MODEL_NAME)

def preprocess_function(examples):
    model_inputs = tokenizer(examples["input_text"], max_length=128, truncation=True, padding="max_length")
    labels = tokenizer(examples["target_text"], max_length=128, truncation=True, padding="max_length")

    labels["input_ids"] = [
        [1 if l != tokenizer.pad_token_id else -100 for l in label] for label in labels["input_ids"]
    ]

    model_inputs["labels"] = labels["input_ids"]
    return model_inputs

tokenized_datasets = dataset.map(preprocess_function, batched=True)

/usr/local/lib/python3.12/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
  no secret 'HF_TOKEN' does not exist in your Colab secrets.
  > authenticate with the Hugging Face Hub, create a token in your settings tab (https://huggingface.co/settings/tokens), set it as secret in your Google Colab and restart your session.
  you will be able to reuse this secret in all of your notebooks.
  > note that authentication is recommended but still optional to access public models or datasets.
  warnings.warn(
  tokenizer_config.json: 2.54kB? [00:00<00:00, 142kB/s]
  piece model: 100% 792k/792k [00:01<00:00, 601kB/s]
```

Disk 72.86 GB available

Load Model with LoRA.ipynb

Insert Runtime Tools Help

+ Text Run all

Load Model with LoRA

```
[1]
# Load base model
base_model = AutoModelForSeq2SeqLM.from_pretrained(MODEL_NAME, torch_dtype=torch.float32)

# Define LoRA
lora_config = LoraConfig(
    r=16,
    lora_alpha=32,
    target_modules=["q", "v"],
    lora_dropout=0.05,
    bias="none",
    task_type=TaskType.SEQ_2_SEQ_LM
)

config.json: 1.40kB? [00:00<00:00, 120kB/s]
`torch_dtype` is deprecated! Use `dtype` instead!
model.safetensors: 100% 990M/990M [00:11<00:00, 51.4MB/s]
generation_config.json: 100% 147/147 [00:00<00:00, 10.7kB/s]
```

Apply LoRA to the model

```
[1]
model = get_peft_model(base_model, lora_config)
model.print_trainable_parameters()
```

trainable params: 1,769,472 || all params: 249,347,328 || trainable%: 0.709%

Training Arguments

74.14 GB available

LoRA.ipynb

Insert Runtime Tools Help

+ Text Run all

Training Arguments

```
[ ] training_args = Seq2SeqTrainingArguments(  
    output_dir=OUTPUT_DIR,  
    per_device_train_batch_size=BATCH_SIZE,  
    per_device_eval_batch_size=BATCH_SIZE,  
    learning_rate=3e-4,  
    num_train_epochs=5,  
    logging_dir=f'{OUTPUT_DIR}/logs',  
    logging_strategy="steps",  
    logging_steps=20,  
    eval_strategy="epoch",  
    save_strategy="epoch",  
    save_total_limit=2,  
    predict_with_generate=True,  
    fp16=False,  
    push_to_hub=False,  
)
```

Trainer

```
[ ] trainer = Seq2SeqTrainer(  
    model=model,  
    args=training_args,  
    train_dataset=tokenized_datasets["train"],  
    eval_dataset=tokenized_datasets["test"],  
    tokenizer=tokenizer,  
    data_collator=DataCollatorForSeq2Seq(tokenizer, model=model)  
)  
  
print("Starting training...")  
trainer.train()
```

74.14 GB available

LoRA.ipynb

Insert Runtime Tools Help

+ Text Run all

Syncing run **treasured-aardvark-2** to Weights & Biases (docs)
View project at <https://wandb.ai/anujanaru812/university-of-moratuwa/huggingface>
View run at <https://wandb.ai/anujanaru812/university-of-moratuwa/huggingface/runs/us8abugs> [361/450 02:37 < 00:39, 2.28 it/s, Epoch 4/5]

Epoch	Training Loss	Validation Loss
1	2.046800	1.496319
2	1.090500	0.661210
3	0.794400	0.411531

[9/10 00:01 < 00:00, 4.94 it/s]
[450/450 03:20, Epoch 5/5]

Epoch	Training Loss	Validation Loss
1	2.046800	1.496319
2	1.090500	0.661210
3	0.794400	0.411531
4	0.627800	0.346995
5	0.578300	0.324698

```
TrainOutput(global_step=450, training_loss=1.207777435514662, metrics={'train_runtime': 565.9482, 'train_samples_per_second': 6.361, 'train_steps_per_second': 0.795, 62117385338880.0, 'train_loss': 1.207777435514662, 'epoch': 5.0})
```

```
[ ] model.save_pretrained(OUTPUT_DIR)  
tokenizer.save_pretrained(OUTPUT_DIR)  
  
['./banking_assistant_lora/tokenizer_config.json',  
 './banking_assistant_lora/special_tokens_map.json',  
 './banking_assistant_lora/spiece.model',  
 './banking_assistant_lora/added_tokens.json',  
 './banking_assistant_lora/tokenizer.json')
```

LoRA.ipynb

Insert Runtime Tools Help

+ Text Run all

Compare Base vs Fine-Tuned

```
[ ] import pandas as pd  
from transformers import pipeline  
  
print("\n--- Generating Comparisons (Base vs Fine-Tuned) ---")  
  
test_prompts = [  
    "I lost my credit card, help!",  
    "What is the balance of my Savings account?",  
    "Transfer 500 to Mom.",  
    "When does the bank close?",  
    "Can I get a home loan?",  
    "My platinum card was stolen.",  
    "Send 200.50 to Alice.",  
    "Do you are open on Sundays?",  
    "What is the interest rate for a car loan?",  
    "I need to block my debit card."  
]  
  
def generate_response(model, tokenizer, prompt):  
    # Standard T5 format  
    input_text = f"Instruction: {prompt}"  
    inputs = tokenizer(input_text, return_tensors="pt").to(model.device)  
  
    # Generate  
    outputs = model.generate(**inputs, max_new_tokens=64)  
    return tokenizer.decode(outputs[0], skip_special_tokens=True)
```

bRA.ipynb

File Runtime Tools Help

+ Text ▶ Run all ▾

Evaluation & Visualization

Plot training loss

```
[ ] import matplotlib.pyplot as plt
import pandas as pd
```

```
[ ] history = trainer.state.log_history
loss_data = {'step': [], 'loss': []}

for entry in history:
    if 'loss' in entry:
        loss_data['step'].append(entry['step'])
        loss_data['loss'].append(entry['loss'])

plt.figure(figsize=(10, 5))
plt.plot(loss_data['step'], loss_data['loss'], label='Training Loss')
plt.xlabel('Step')
plt.ylabel('Loss')
plt.title('Training Loss Curve (LoRA Fine-tuning)')
plt.legend()
plt.grid(True)
plt.savefig("training_loss.png")
plt.show()
```

bRA.ipynb

File Runtime Tools Help

+ Text ▶ Run all ▾

```
[ ] results = []

print(f"[{'Instruction':<48} | {'BASE Model (Before)':<30} | {'FINE-TUNED (After)'})"
print("-" * 110)

for prompt in test_prompts:
    # 1. Get Fine-Tuned Response
    ft_response = generate_response(model, tokenizer, prompt)

    # 2. Get Base Model Response (Temporarily disable LoRA)
    with model.disable_adapter():
        base_response = generate_response(model, tokenizer, prompt)

    print(f"(prompt:<48) | {base_response:<30} | {ft_response})")

    results.append({
        "instruction": prompt,
        "base_model": base_response,
        "fine_tuned_model": ft_response
    })

# Save comparisons to CSV for your submission
df = pd.DataFrame(results)
df.to_csv("model_comparison.csv", index=False)
print("\nSuccess! Saved comparison results to 'model_comparison.csv'")
```

--- Generating Comparisons (Base vs Fine-Tuned) ---

Instruction	BASE Model (Before)	FINE-TUNED (After)
I lost my credit card, help!	I lost my credit card, help!	I have locked your credit card temporarily. Please check your mobile app to permanently block it or request a new one.
What is the balance of my Savings account? 0	The balance of your Savings is \$2682.00.	
Transfer 500 to Mom.	500 to Mom.	Transfer request for 500 received. Please verify the recipient's details before I confirm.
When does the bank close?	Late at night	We are currently closed. Our hours are 9 AM to 5 PM, Monday through Friday.
Can I get a home loan?	Yes	You can apply for a Home Loan directly through the app under the 'Loans' tab.
My platinum card was stolen.	I'm sorry.	Please log in to your Apple app and select 'Pepper Card' immediately. I can also connect you to a secure line for further assistance.
Send 200.50 to Alice.	Alice: 200.50.	Please send 200.50 to Alice. Please send it to your phone.
Do you open on Sundays?	No	We are closed on Sundays and public holidays.
What is the interest rate for a car loan? 0		You can apply for a car loan with a fixed interest rate. However, you can apply for a car loan with a flexible interest rate.
I need to block my debit card.	To whom a you to do this?	In block your debit card, you need to enter your DTM (that we never had).

74.14 GB available